

SAT with Memory

Md Solimul Chowdhury


Automated Reasoning Group
Amazon Web Services

chowmdso@amazon.com

September/24/2021

Boolean Satisfiability (SAT) Solving

- SAT solvers are well-known engines of proof generation
- SAT task \longrightarrow to **prove Satisfiability** of a Boolean formula \mathcal{F}
 - Finding an assignment to the variables such that \mathcal{F} is *True* (**SAT**)
 - Otherwise, report Unsatisfiability (**UNSAT**)



- important use cases in AWS
- e.g., verification of C programs via model checking

SAT Meets the Cloud

- **Lambda Service**

- serverless | | program executables → **convenience**

Use Case 1: SAT as a Lambda Service.



? Solve this SAT
Problem

Lambda Service

: SAT/UNSAT


: Sorry! Timeout (15 minutes)

Issue: Solves 99% within time out. The other 1%?

- In cloud, we require **migration** of SAT jobs!

SAT Meets the Cloud

- Incremental Solving is in many applications of SAT
 - given a SAT problem \mathbf{P} and a set of constraints $\{C_1, \dots, C_n\}$,
 - n incremental calls: $\text{solve}(\mathbf{P})|_{C_1} \dots \text{solve}(\mathbf{P})|_{C_n}$
 - For an increment, no need to solve \mathbf{P} from the scratch
- Example: Reachable paths to a node A from X
 - $\text{solve}(X, A)|_{C_1} \dots \text{solve}(X, A)|_{C_n}$



- In cloud, we prefer **session free** incremental SAT solving!

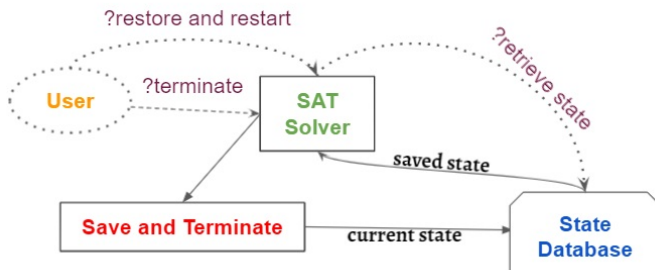
Our Solution : Saving and Reusing SAT States

Issue 1	Issue 2
Lacking of Migration Support for SAT jobs	Requirement for Session Maintenance for Incremental SAT jobs

- To solve these issues, we propose
 - to save states at the end of a SAT call given a timeout
 - and reuse them in the next call
- Our solution,
 - preserve provability,
 - has low overhead, and
 - has a low storage requirement

The Concept

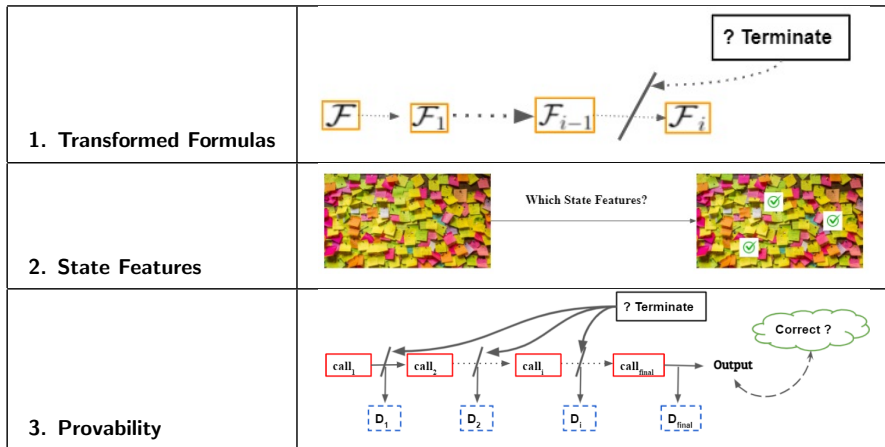
- Stop (at request), Save State, Solve (at request)



- We propose SAT_{mem} , *SAT with Memory*

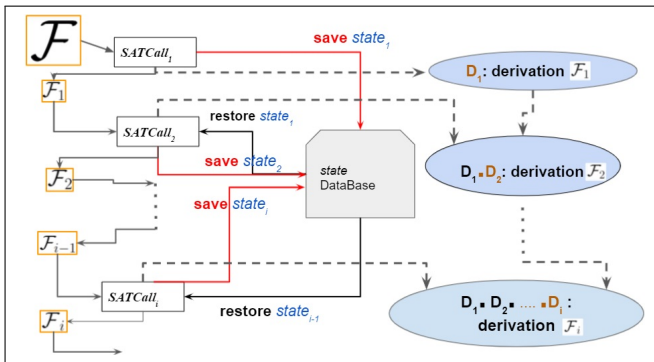
eifjccikedgrdlguklhjknntudeghrkkrburtftdntf
eifjccikedgrcfdhjbeijjcdkvgrbknkbltvjgvgflh

Design Options for SAT_{mem}



The SAT_{mem} Architecture

Save and Reuse	Derivation Logging
A. States B. Transformed Input Formula - Solver changes \mathcal{F}	A. Log Derivations by concatenation. B. Verification of combined derivation. - Debugging. - UNSAT verification. - Satisfiability Assignment Extraction for \mathcal{F} .



Implementation and Benchmarks

- **Implementation:**

- SAT_{mem} on top of CaDiCaL (2021)

- Top performing SAT solver + supports incremental solving

- CaDiCaL \rightarrow CaDiCaL_{mem}

- Number of calls CaDiCaL
 - each call feeds itself state saved by the previous call

- **Benchmarks:** 297 instances (SR19) from SAT Race-2019

- CaDiCaL runs more than 300 seconds for each of these instances
 - In CaDiCaL_{mem}, we use 300 secs as timeout threshold

Experimental Setup

- CaDiCaL: 297 SR19 instances with a timeout of 2000 seconds

CaDiCaL	
Call1	
2000 secs	

- CaDiCaL_{mem}: Same 297 instances, but two different modes of termination

CaDiCaL _{mem} Mode 1	
Call1	Call2
300 secs	1700 secs

CaDiCaL _{mem} Mode 2						
Call1	Call2	Call3	Call4	Call5	Call6	Call7
300 secs	300 secs	300 secs	300 secs	300 secs	300 secs	200 secs

Experimental Setup



- At termination, the current call saves the following state features:
 - 1 Transformed Formula
 - 2 Internal to External Map / External to Internal Map
 - 3 Heuristic score of variables
 - 4 Phases
 - 5 Learned Clauses
- Details will be available in: https://w.amazon.com/bin/view/ARG/ProofPlatformsTeam/SAT_with_Memory

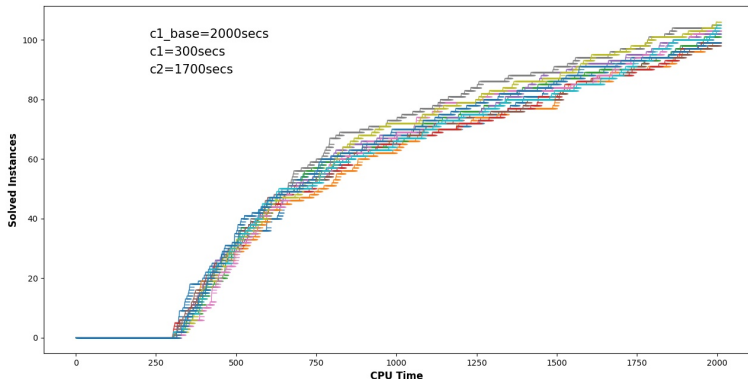
Run-time Variation with Termination by Wall Clock

- In two different runs, at termination, a **call** may end into two different states This induces
 - non-determinism in the **next call**
 - and runtime variation
- For experiments, we needed to perform multiple runs to estimate the average case ...

CaDiCaL_{mem} Mode 1 results

CaDiCaL _{mem} Mode 1	
Call1	Call2
300 secs	1700 secs

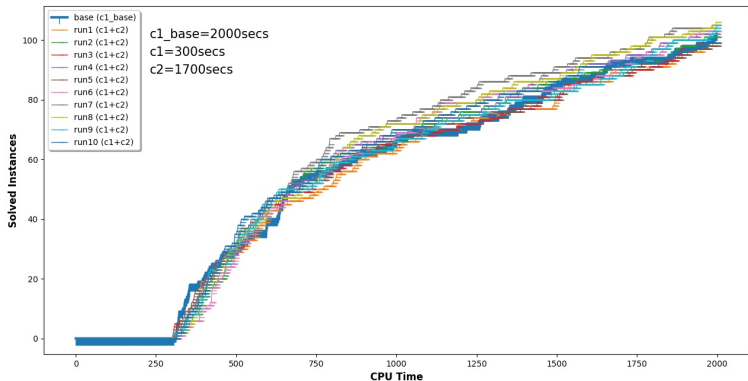
- Compared 10 runs of CaDiCaL_{mem} (Mode 1) with CaDiCaL (base)
 - Which line is the base? Any guess?



CaDiCaL_{mem} Mode 1 results : Low overhead

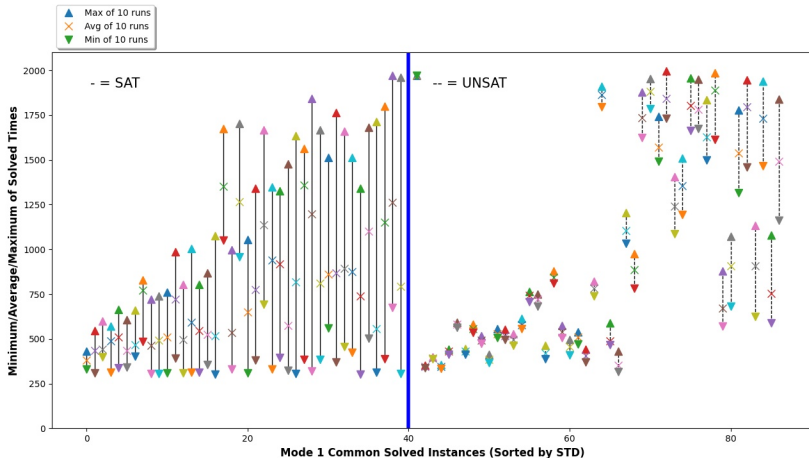
- Compared 10 runs of CaDiCaL_{mem} (Mode 1) with CaDiCaL (base)

CaDiCaL _{mem} Mode 1	
Call1	Call2
300 secs	1700 secs



Larger Runtime Variation with SAT instances (Mode 1)

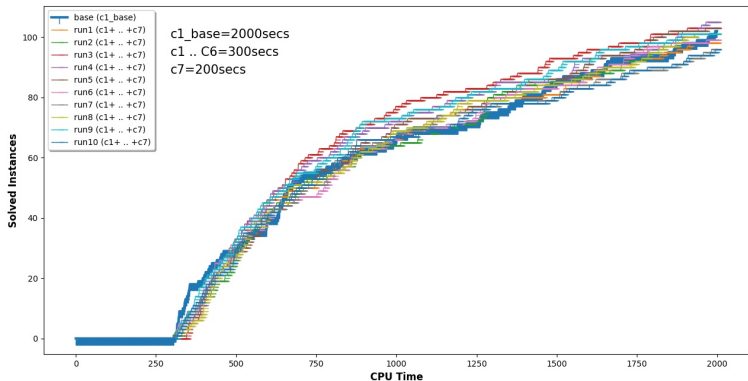
- Statistics on Common solved instances (86) over 10 runs



CaDiCaL_{mem} Mode 2 results: Low overhead

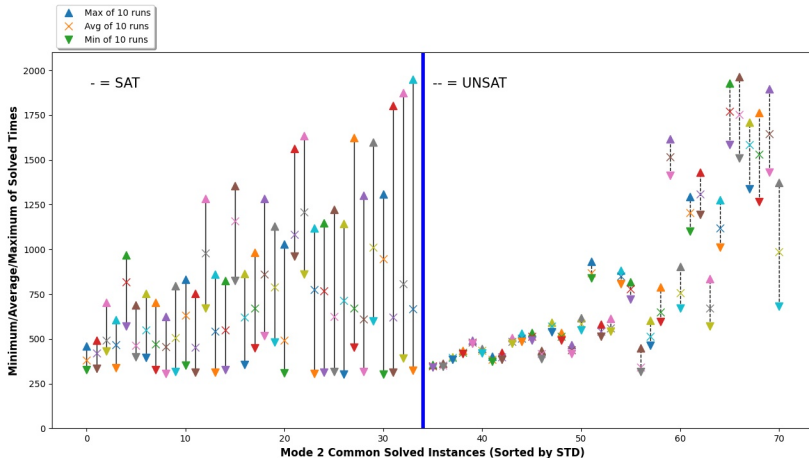
- Compared 10 runs of CaDiCaL_{mem} (Mode 2) with CaDiCaL (base)

CaDiCaL _{mem} Mode 2						
Call1	Call2	Call3	Call4	Call5	Call6	Call7
300 secs	300 secs	300 secs	300 secs	300 secs	300 secs	200 secs



Larger Runtime Variation with SAT Instances (Mode 2)


- Statistics on Common solved instances (70) over 10 runs.




Low Storage Requirement with CaDiCaL_{mem}

- Single run of Mode 1
- Peak Memory with Call 1
- State DB in text file (no compression)

Comparison: Size of State DB and Formula



Comparison: Size of State DB and Peak Memory



Summary and Future Steps

• Summary

- Many shot SAT solving with memory. For AWS
 - Migration of SAT Jobs
 - Session free Incremental SAT jobs
 - Restart SAT jobs
- Small overhead
- Light storage requirement
- Proof producing \longrightarrow verifiable

• Whats Next?

- Experiments with few more sets of SAT competition benchmarks
- Experiments focusing the runtime variations
 - Exploitable in parallel settings?
- Experiments with incremental case
- Constructions of Transformed formula and Learned Clauses from derivation
 - Solver independent implementation

Acknowledgement

- Marijn Heule (Mentor)
- Mike Whalen (Manager)
- Frankie Botero